Analytics Vidhya

# LLM API Pricing Awareness


Instructor

Dipanjan Sarkar

Head of Community & Principal AI Scientist at Analytics Vidhya

Google Developer Expert - ML & Cloud Champion Innovator

Published Author

# LLMs are usually not free and have a cost

- Commercial LLMs typically charge based on the number of tokens or requests made.
- Open-source LLMs typically incur costs for maintaining GPU infrastructure or based on tokens or requests if using a managed platform.

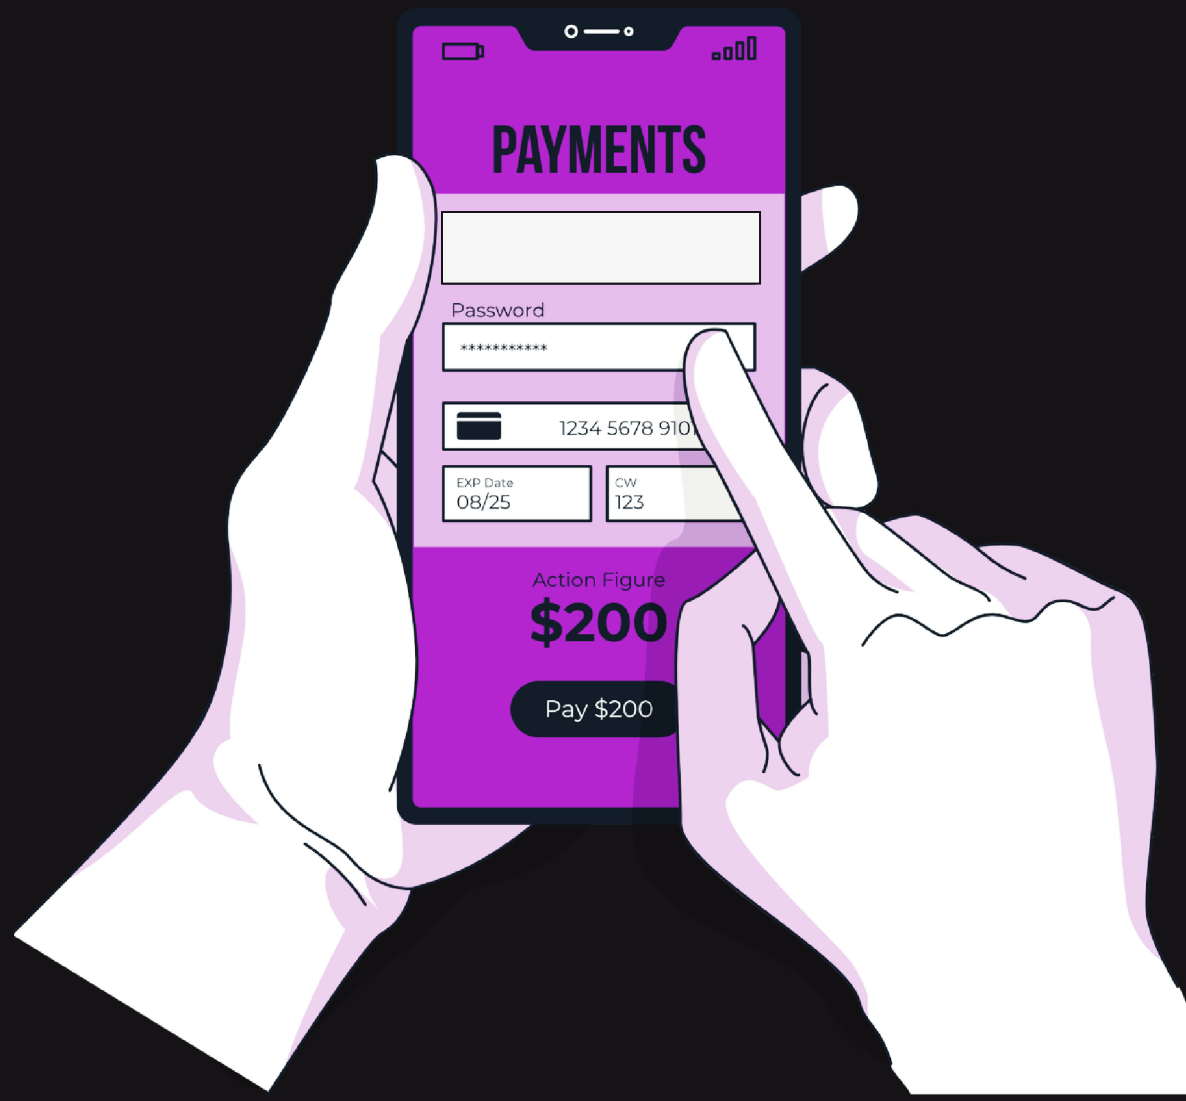

*It's Good to be aware of API costs to prevent unnecessary surprises when using them...*

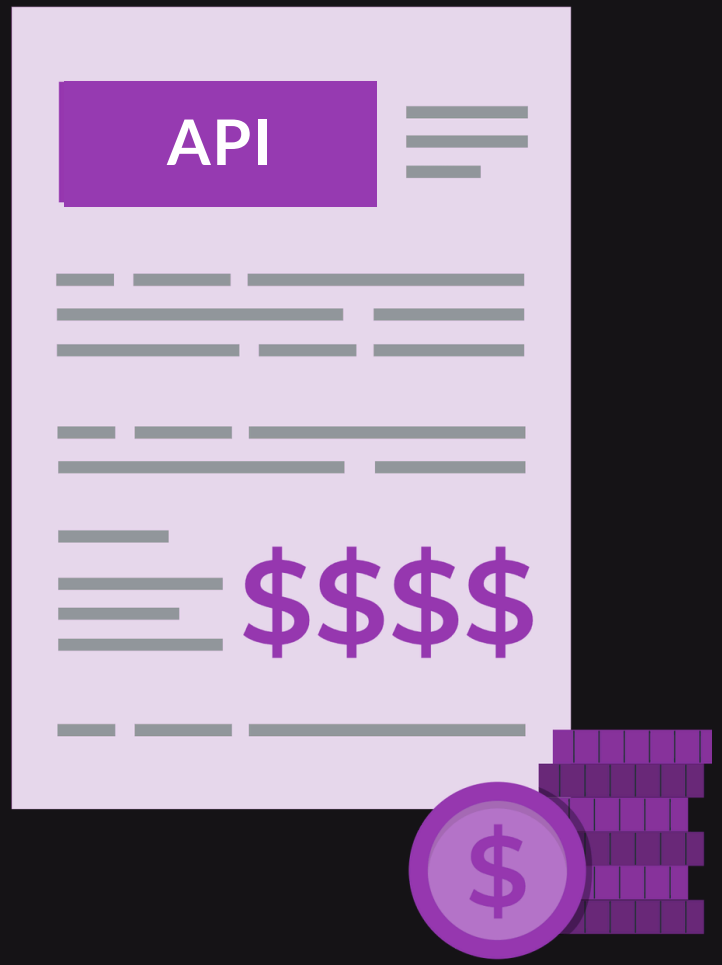

# Commercial LLM APIs

OpenAI GPT API

Google Gemini API

# OpenAI API Pricing (As of Late 2024)

| Model | Input Pricing ($ per 1K tokens) | Output Pricing ($ per 1K tokens) |
|---|---|---|
| gpt-4o | 0.00250 | 0.0100 |
| gpt-4o-mini | 0.00015 | 0.0006 |
| gpt-3.5-turbo | 0.00050 | 0.0015 |
| o1-preview | 0.01500 | 0.0600 |
| o1-mini | 0.00300 | 0.0120 |
| chatgpt-4o-latest | 0.00500 | 0.0150 |

Source: https://openai.com/api/pricing/ for the latest prices

# Google Gemini API Pricing (As of Late 2024)

| Model | Rate Limits | Input Pricing ($ per 1K tokens) | Input Pricing ($ per 1K tokens) |
|---|---|---|---|
| Gemini 1.5 Flash Free | 15 RPM, 1M TPM, 1,500 RPD | Free of Charge | Free of Charge |
| Gemini 1.5 Flash Paid ( ≤ 128k tokens ) | 2,000 RPM, 4M TPM | 0.000075 | 0.00030 |
| Gemini 1.5 Flash Paid ( > 128k tokens ) | 2,000 RPM, 4M TPM | 0.00015 | 0.00060 |
| Gemini 1.5 Pro Free | 2 RPM, 32k TPM, 50 RPD | Free of Charge | Free of Charge |
| Gemini 1.5 Pro Paid ( ≤ 128k tokens ) | 1,000 RPM, 4M TPM | 0.00125 | 0.00500 |
| Gemini 1.5 Pro Paid ( > 128k tokens ) | 1,000 RPM, 4M TPM | 0.00250 | 0.01000 |

**Source: https://ai.google.dev/pricing for the latest prices**

# Open-Source LLM API Platforms

Groq

HuggingFace

# HuggingFace APIs

- The **Serverless Inference API** provides fast, free access to thousands of ML models for various tasks, ideal for prototyping and experimentation:
  - **Text Generation**: High-quality responses with large language models.
  - **Image Generation**: Create custom images, including LoRAs for unique styles.
  - **Document Embeddings**: Build advanced search and retrieval systems.
  - **Classical AI Tasks**: Models for classification, speech recognition, and more.

# HuggingFace APIs

## ⚡ Fast and Free to Get Started:

- The Inference API is free with higher rate limits for PRO users.
- For production needs, explore **Inference Endpoints** for dedicated resources, autoscaling, advanced security features, and more.

# Groq Cloud Free Tier Limits (Late 2024)

| ID | Requests per Minute | Requests per Day | Tokens per Minute | Tokens per Day |
|---|---|---|---|---|
| gemma-7b-it | 30 | 14400 | 15,000 | 500,000 |
| gemma2-9b-it | 30 | 14400 | 15,000 | 500,000 |
| llama-3.1-70b-versatile | 30 | 14400 | 18,000 | 500,000 |
| llama-3.1-8b-instant | 30 | 14400 | 20,000 | 500,000 |
| llama-3.2-11b-text-preview | 30 | 7000 | 7,000 | 500,000 |
| llama-3.2-11b-vision-preview | 30 | 7000 | 7,000 | 500,000 |
| llama-3.2-1b-preview | 30 | 7000 | 7,000 | 500,000 |
| llama-3.2-2b-preview | 30 | 7000 | 7,000 | 500,000 |
| llama-3.2-90b-text-preview | 30 | 7000 | 7,000 | 500,000 |
| llama-3.2-90b-vision-preview | 30 | 7000 | 7,000 | 500,000 |
| llama-guard-3-8b | 15 | 3500 | 7,000 | 250,000 |
| llama3-70b-8192 | 30 | 14400 | 15,000 | 500,000 |
| llama3-8b-8192 | 30 | 14400 | 6,000 | 500,000 |
| llama3-groq-70b-8192-tool-preview | 30 | 14400 | 30,000 | 500,000 |
| llama3-groq-8b-8192-tool-preview | 30 | 14400 | 30,000 | 500,000 |
| llava-v1.5-7b-4096-preview | 30 | 14400 | 15,000 | 500,000 |
| mixtral-8x7b-32768 | 30 | 14400 | 5,000 | 500,000 |

Check out pricing details here for free API and here for paid API

Analytics Vidhya

# Use Claude to Improve Existing Prompts Faster

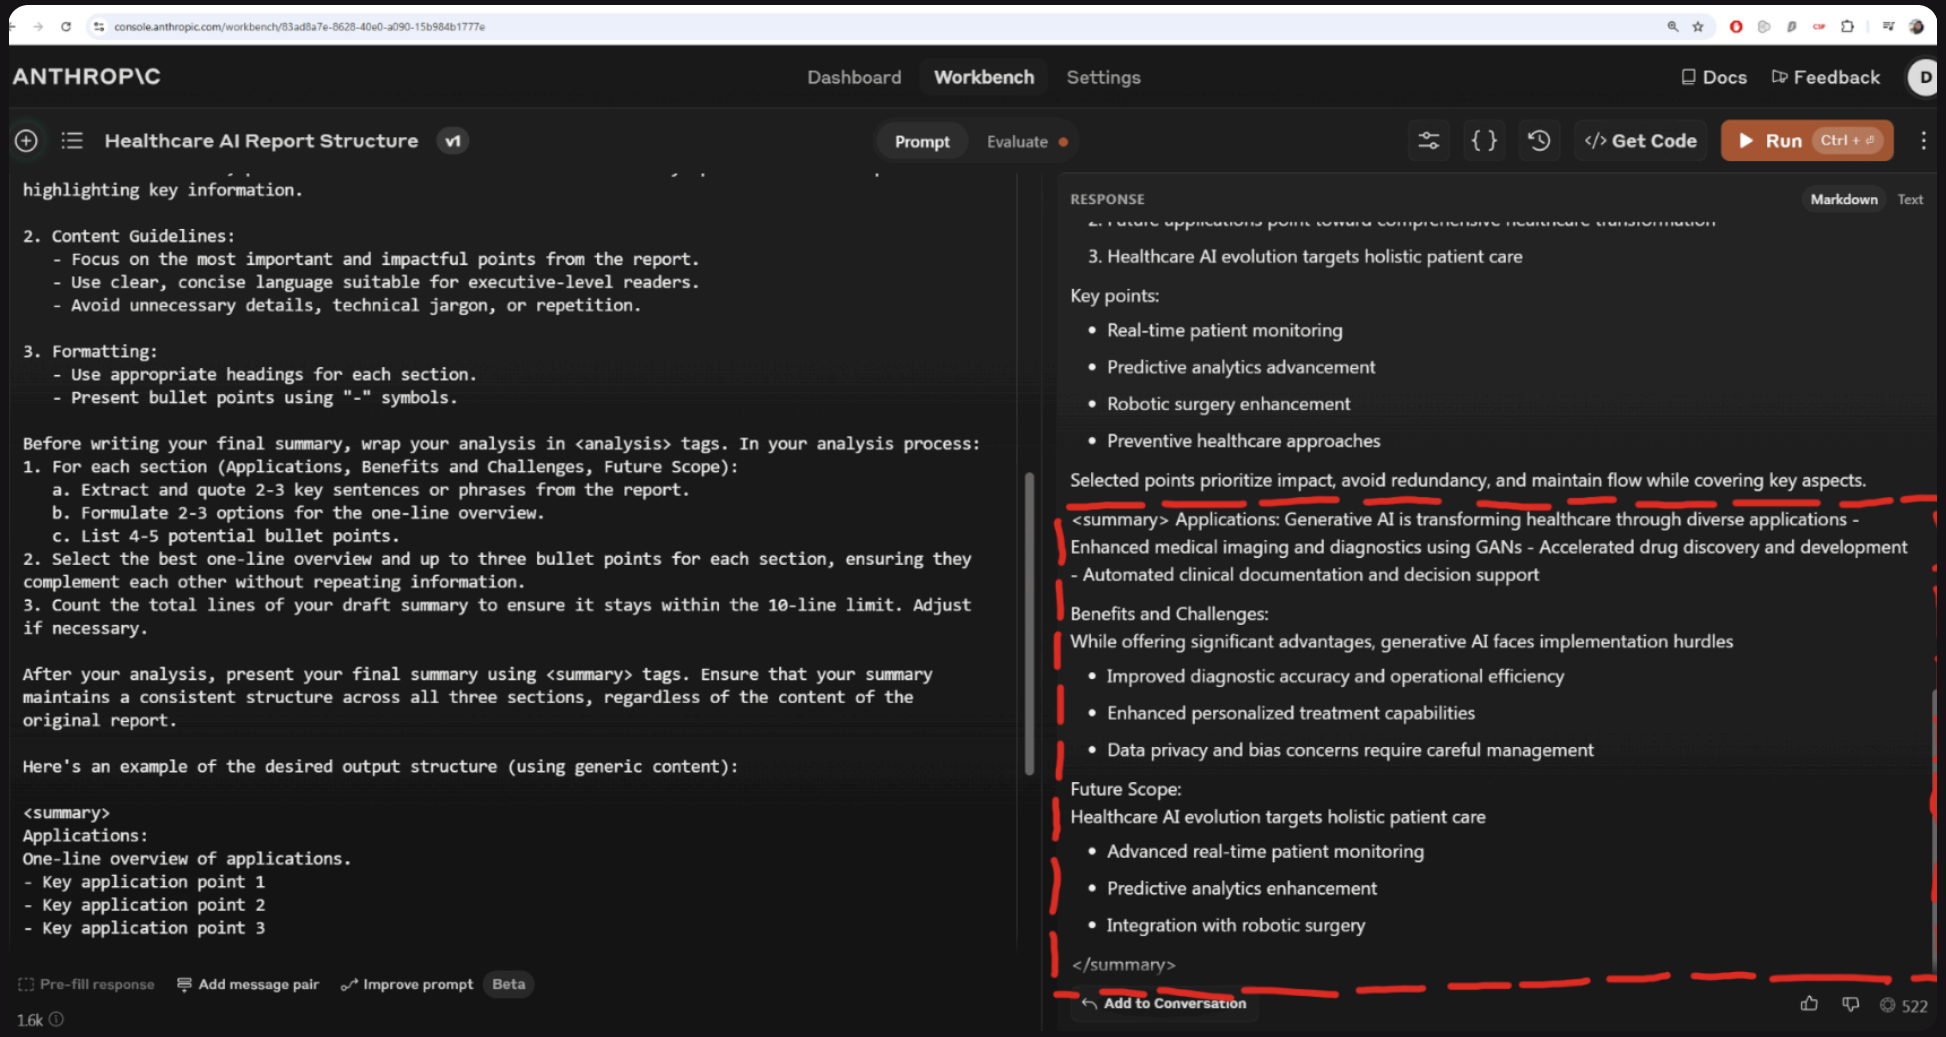


Analytics Vidhya

# Thank You

Analytics Vidhya